

# Assignment4\_ANN

October 27, 2024

## 1 Semiconductor manufacturing process dataset

### 1.1 Project Description

Source: <https://www.kaggle.com/saurabhbhagchi/fmst-semiconductor-manufacturing-project>

A complex modern semiconductor manufacturing process is normally under constant surveillance via the monitoring of signals/variables collected from sensors and or process measurement points. However, not all of these signals are equally valuable in a specific monitoring system. The measured signals contain a combination of useful information, irrelevant information as well as noise. Engineers typically have a much larger number of signals than are actually required. If we consider each type of signal as a feature, then feature selection may be applied to identify the most relevant signals. The Process Engineers may then use these signals to determine key factors contributing to yield excursions downstream in the process. This will enable an increase in process throughput, decreased time to learning, and reduce per-unit production costs. These signals can be used as features to predict the yield type. And by analyzing and trying out different combinations of features, essential signals that are impacting the yield type can be identified.

Dataset: SemiconductorManufacturingProcessDataset.csv (on Canvas)

Later, we will learn how to apply PCA (Principal Component Analyses) for feature selection; then we will apply ANN to predict the Pass/Fail. in this exercise our objective is to repeat the same steps we did above for Supplier Data: Cleaning & Scaling Data, Encode Categorical Data, Split the Data to Training & Test Sets.

### 1.2 Importing the Libraries

```
[188]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

### 1.3 Importing the Dataset

```
[189]: dataset = pd.read_csv('SemiconductorManufacturingProcessDataset.csv')
```

## 1.4 Showing the Dataset in a Table

```
[190]: pd.DataFrame(dataset)
       #dataset
```

```
[190]:
```

		Time	Sensor 1	Sensor 2	Sensor 3	Sensor 4	Sensor 5	\
0	7/19/2008	11:55	3030.93	2564.00	2187.7333	1411.1265	1.3602	
1	7/19/2008	12:32	3095.78	2465.14	2230.4222	1463.6606	0.8294	
2	7/19/2008	13:17	2932.61	2559.94	2186.4111	1698.0172	1.5102	
3	7/19/2008	14:43	2988.72	2479.90	2199.0333	909.7926	1.3204	
4	7/19/2008	15:22	3032.24	2502.87	2233.3667	1326.5200	1.5334	
...	...	...	...	...	...	...	...	
1562	10/16/2008	15:13	2899.41	2464.36	2179.7333	3085.3781	1.4843	
1563	10/16/2008	20:49	3052.31	2522.55	2198.5667	1124.6595	0.8763	
1564	10/17/2008	5:26	2978.81	2379.78	2206.3000	1110.4967	0.8236	
1565	10/17/2008	6:01	2894.92	2532.01	2177.0333	1183.7287	1.5726	
1566	10/17/2008	6:07	2944.92	2450.76	2195.4444	2914.1792	1.5978	

		Sensor 6	Sensor 7	Sensor 8	Sensor 9	...	Sensor 429	Sensor 430	\
0		97.6133	0.1242	1.5005	0.0162	...	14.9509	0.5005	
1		102.3433	0.1247	1.4966	-0.0005	...	10.9003	0.5019	
2		95.4878	0.1241	1.4436	0.0041	...	9.2721	0.4958	
3		104.2367	0.1217	1.4882	-0.0124	...	8.5831	0.4990	
4		100.3967	0.1235	1.5031	-0.0031	...	10.9698	0.4800	
...		...	...	...	...	...	...	...	
1562		82.2467	0.1248	1.3424	-0.0045	...	11.7256	0.4988	
1563		98.4689	0.1205	1.4333	-0.0061	...	17.8379	0.4975	
1564		99.4122	0.1208	NaN	NaN	...	17.7267	0.4987	
1565		98.7978	0.1213	1.4622	-0.0072	...	19.2104	0.5004	
1566		85.1011	0.1235	NaN	NaN	...	22.9183	0.4987	

		Sensor 431	Sensor 432	Sensor 433	Sensor 434	Sensor 435	Sensor 436	\
0		0.0118	0.0035	2.3630	NaN	NaN	NaN	
1		0.0223	0.0055	4.4447	0.0096	0.0201	0.0060	
2		0.0157	0.0039	3.1745	0.0584	0.0484	0.0148	
3		0.0103	0.0025	2.0544	0.0202	0.0149	0.0044	
4		0.4766	0.1045	99.3032	0.0202	0.0149	0.0044	
...		...	...	...	...	...	...	
1562		0.0143	0.0039	2.8669	0.0068	0.0138	0.0047	
1563		0.0131	0.0036	2.6238	0.0068	0.0138	0.0047	
1564		0.0153	0.0041	3.0590	0.0197	0.0086	0.0025	
1565		0.0178	0.0038	3.5662	0.0262	0.0245	0.0075	
1566		0.0181	0.0040	3.6275	0.0117	0.0162	0.0045	

		Sensor 437	Pass/Fail
0		NaN	Pass
1		208.2045	Pass

2	82.8602	Fail
3	73.8432	Pass
4	73.8432	Pass
...	...	...
1562	203.1720	Pass
1563	203.1720	Pass
1564	43.5231	Pass
1565	93.4941	Pass
1566	137.7844	Pass

[1567 rows x 439 columns]

## 1.5 A Quick Review of the Data

```
[191]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1567 entries, 0 to 1566
Columns: 439 entries, Time to Pass/Fail
dtypes: float64(437), object(2)
memory usage: 5.2+ MB
```

## 1.6 Seperate The Input and Output

Here, we put the independent variables in X and the dependent variable in y.

```
[192]: X = dataset.iloc[:, 1:438].values
       y = dataset.iloc[:, -1].values
```

## 1.7 Showing the Input Data in a Table format

```
[193]: pd.DataFrame(X)
```

```
[193]:
```

	0	1	2	3	4	5	6	\		
0	3030.93	2564.00	2187.7333	1411.1265	1.3602	97.6133	0.1242			
1	3095.78	2465.14	2230.4222	1463.6606	0.8294	102.3433	0.1247			
2	2932.61	2559.94	2186.4111	1698.0172	1.5102	95.4878	0.1241			
3	2988.72	2479.90	2199.0333	909.7926	1.3204	104.2367	0.1217			
4	3032.24	2502.87	2233.3667	1326.5200	1.5334	100.3967	0.1235			
...	...	...	...	...	...	...	...			
1562	2899.41	2464.36	2179.7333	3085.3781	1.4843	82.2467	0.1248			
1563	3052.31	2522.55	2198.5667	1124.6595	0.8763	98.4689	0.1205			
1564	2978.81	2379.78	2206.3000	1110.4967	0.8236	99.4122	0.1208			
1565	2894.92	2532.01	2177.0333	1183.7287	1.5726	98.7978	0.1213			
1566	2944.92	2450.76	2195.4444	2914.1792	1.5978	85.1011	0.1235			
	7	8	9	...	427	428	429	430	431	\
0	1.5005	0.0162	-0.0034	...	1.6765	14.9509	0.5005	0.0118	0.0035	

1	1.4966	-0.0005	-0.0148	...	1.1065	10.9003	0.5019	0.0223	0.0055
2	1.4436	0.0041	0.0013	...	2.0952	9.2721	0.4958	0.0157	0.0039
3	1.4882	-0.0124	-0.0033	...	1.7585	8.5831	0.4990	0.0103	0.0025
4	1.5031	-0.0031	-0.0072	...	1.6597	10.9698	0.4800	0.4766	0.1045
...	...	...	...	...	...	...	...	...	...
1562	1.3424	-0.0045	-0.0057	...	1.4879	11.7256	0.4988	0.0143	0.0039
1563	1.4333	-0.0061	-0.0093	...	1.0187	17.8379	0.4975	0.0131	0.0036
1564	NaN	NaN	NaN	...	1.2237	17.7267	0.4987	0.0153	0.0041
1565	1.4622	-0.0072	0.0032	...	1.7085	19.2104	0.5004	0.0178	0.0038
1566	NaN	NaN	NaN	...	1.2878	22.9183	0.4987	0.0181	0.0040

	432	433	434	435	436
0	2.3630	NaN	NaN	NaN	NaN
1	4.4447	0.0096	0.0201	0.0060	208.2045
2	3.1745	0.0584	0.0484	0.0148	82.8602
3	2.0544	0.0202	0.0149	0.0044	73.8432
4	99.3032	0.0202	0.0149	0.0044	73.8432
...	...	...	...	...	...
1562	2.8669	0.0068	0.0138	0.0047	203.1720
1563	2.6238	0.0068	0.0138	0.0047	203.1720
1564	3.0590	0.0197	0.0086	0.0025	43.5231
1565	3.5662	0.0262	0.0245	0.0075	93.4941
1566	3.6275	0.0117	0.0162	0.0045	137.7844

[1567 rows x 437 columns]

## 1.8 A Quick Check of the Output Data

```
[194]: pd.DataFrame(y)
```

```
[194]:
      0
0    Pass
1    Pass
2    Fail
3    Pass
4    Pass
...
1562 Pass
1563 Pass
1564 Pass
1565 Pass
1566 Pass
```

[1567 rows x 1 columns]

## 1.9 Taking care of missing data

```
[195]: from sklearn.impute import SimpleImputer
imputer = SimpleImputer(missing_values=np.nan, strategy='mean')
imputer.fit(X)
X = imputer.transform(X)
```

```
[196]: # A quick check
print(X)
```

```
[[3.03093000e+03 2.56400000e+03 2.18773330e+03 ... 1.64749042e-02
 5.28333333e-03 9.96700663e+01]
 [3.09578000e+03 2.46514000e+03 2.23042220e+03 ... 2.01000000e-02
 6.00000000e-03 2.08204500e+02]
 [2.93261000e+03 2.55994000e+03 2.18641110e+03 ... 4.84000000e-02
 1.48000000e-02 8.28602000e+01]
 ...
 [2.97881000e+03 2.37978000e+03 2.20630000e+03 ... 8.60000000e-03
 2.50000000e-03 4.35231000e+01]
 [2.89492000e+03 2.53201000e+03 2.17703330e+03 ... 2.45000000e-02
 7.50000000e-03 9.34941000e+01]
 [2.94492000e+03 2.45076000e+03 2.19544440e+03 ... 1.62000000e-02
 4.50000000e-03 1.37784400e+02]]
```

## 1.10 Encoding Categorical Data

### 1.10.1 Encoding the Independent Variable

```
[197]: # we don't have any categorical data
```

### 1.10.2 Encoding the Dependent Variable

```
[198]: from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
y = le.fit_transform(y)
```

```
[199]: # a quick check
print(y)
```

```
[1 1 0 ... 1 1 1]
```

## 1.11 Feature Scaling

```
[200]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X = sc.fit_transform(X)
```

### 1.12 Splitting the Dataset into the Training set and Test set

[illegible]

```
[202]: print(X_train)
```

```

[[-0.96046311 -0.73813734 -0.92237938 ... -0.06531583 -0.16868853
  -0.2120265 ]
 [-0.87742151  0.5426257  -0.13250295 ...  0.60499301  0.38972867
  3.17408017]
 [ 0.05645609 -1.51130825  1.47184855 ...  0.04829584  0.04071792
  0.42803032]
 ...
 [-0.55464836 -0.10473817 -1.25868434 ...  1.9342495   2.16968349
  0.21655552]
 [-0.24467179 -0.00336937 -1.42475658 ... -0.64473532 -0.79690788
  -0.64548212]
 [-0.36283589 -0.07880372  0.55448143 ... -0.19028866 -0.02908423
  1.62346601]]

```

```
[203]: print(X_test)
```

```
[[ [ 0.38726113 -0.80185131 0.15916359 ... -0.66745766 -0.65730358
    -0.56211928]
  [-0.02413509 -0.94137368 -0.96134565 ... -1.08782082 -0.86671003
    -0.69795783]
  [ 0.70418053 0.53813704 -0.68064891 ... -0.06531583 0.00581685
    -0.30543175]
  ...
  [ 0.08504418 0.83962507 -1.23522968 ... -1.12190432 -1.07611647
    -0.7688942 ]
  [ 2.27530028 0.21096393 -0.97685385 ... -0.59929066 -0.55260035
    -0.6417691 ]
  [-0.8523729 -0.23341296 -0.28457524 ... -0.50840133 -0.51769928
    -0.33965675]]]
```

```
[204]: print(y_train)
```

$$[0 \ 1 \ 1 \ \dots \ 1 \ 1 \ 1]$$

```
[205]: print(y_test)
```

[illegible]

```

1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
1 1 1 1 0 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 0]

```

```
[206]: print(X_train)
```

```

[[-0.96046311 -0.73813734 -0.92237938 ... -0.06531583 -0.16868853
  -0.2120265 ]
 [-0.87742151  0.5426257  -0.13250295 ...  0.60499301  0.38972867
   3.17408017]
 [ 0.05645609 -1.51130825  1.47184855 ...  0.04829584  0.04071792
   0.42803032]
 ...
 [-0.55464836 -0.10473817 -1.25868434 ...  1.9342495   2.16968349
   0.21655552]
 [-0.24467179 -0.00336937 -1.42475658 ... -0.64473532 -0.79690788
  -0.64548212]
 [-0.36283589 -0.07880372  0.55448143 ... -0.19028866 -0.02908423
   1.62346601]]

```

```
[207]: print(X_test)
```

```

[[ 0.38726113 -0.80185131  0.15916359 ... -0.66745766 -0.65730358
  -0.56211928]
 [-0.02413509 -0.94137368 -0.96134565 ... -1.08782082 -0.86671003
  -0.69795783]
 [ 0.70418053  0.53813704 -0.68064891 ... -0.06531583  0.00581685
  -0.30543175]
 ...
 [ 0.08504418  0.83962507 -1.23522968 ... -1.12190432 -1.07611647
  -0.7688942 ]
 [ 2.27530028  0.21096393 -0.97685385 ... -0.59929066 -0.55260035
  -0.6417691 ]
 [-0.8523729  -0.23341296 -0.28457524 ... -0.50840133 -0.51769928
  -0.33965675]]

```

## 2 Classification Models

We will investigate a logistic classification model, a Random Forest model and an SVM model. The script below will train each model and output comprehensive data about the performance.

```

[208]: from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix
from sklearn.model_selection import cross_val_score, KFold
import seaborn as sns

```

```

def train_and_evaluate_model(model, X_train, X_test, y_train, y_test,
    ↪ model_name, needs_scaling=False):
    """
    Train and evaluate a single model, including k-fold validation
    """
    # Scale the features if needed (important for SVM)
    if needs_scaling:
        scaler = StandardScaler()
        X_train_scaled = scaler.fit_transform(X_train)
        X_test_scaled = scaler.transform(X_test)
    else:
        X_train_scaled = X_train
        X_test_scaled = X_test

    # Train the model
    model.fit(X_train_scaled, y_train)

    # Make predictions
    y_pred = model.predict(X_test_scaled)
    y_pred_proba = model.predict_proba(X_test_scaled)

    # Calculate metrics
    accuracy = accuracy_score(y_test, y_pred)
    class_report = classification_report(y_test, y_pred, zero_division=0)
    conf_matrix = confusion_matrix(y_test, y_pred)

    # Perform k-fold cross validation
    kf = KFold(n_splits=5, shuffle=True, random_state=42)
    cv_scores = cross_val_score(model, X_train_scaled, y_train, cv=kf,
    ↪ scoring='accuracy')

    # Print results
    print(f"\n{model_name} Results:")
    print("-" * 50)
    print(f"Test Set Accuracy: {accuracy:.4f}")
    print(f"\nK-fold Cross Validation Scores:")
    print(f"Mean CV Accuracy: {cv_scores.mean():.4f} (+/- {cv_scores.std() * 2:.
    ↪ 4f})")
    print(f"Individual fold scores: {[f'{score:.4f}' for score in cv_scores]}")
    print("\nClassification Report:")
    print(class_report)

    # Plot confusion matrix
    plt.figure(figsize=(8, 6))
    sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues')
    plt.title(f'Confusion Matrix - {model_name}')

```



```

plt.ylabel('True Label')
plt.xlabel('Predicted Label')
plt.show()

return {
    'accuracy': accuracy,
    'cv_scores': cv_scores,
    'conf_matrix': conf_matrix,
    'predictions': y_pred,
    'probabilities': y_pred_proba
}

def compare_models(models_results):
    """
    Create comparison visualizations for the models
    """
    # Prepare data for box plot
    model_names = list(models_results.keys())
    cv_scores = [results['cv_scores'] for results in models_results.values()]
    test accuracies = [results['accuracy'] for results in models_results.
↪values()]

    # Create box plot
    plt.figure(figsize=(12, 6))

    # Plot CV scores
    plt.subplot(1, 2, 1)
    plt.boxplot(cv_scores, labels=model_names)
    plt.title('Cross Validation Scores Comparison')
    plt.ylabel('Accuracy')
    plt.xticks(rotation=45)

    # Add individual points for CV scores
    for i in range(len(model_names)):
        x = np.random.normal(i + 1, 0.04, size=len(cv_scores[i]))
        plt.plot(x, cv_scores[i], 'r.', alpha=0.5)

    # Plot test accuracies
    plt.subplot(1, 2, 2)
    plt.bar(model_names, test accuracies)
    plt.title('Test Set Accuracy Comparison')
    plt.ylabel('Accuracy')
    plt.xticks(rotation=45)

    plt.tight_layout()
    plt.show()

```

```

# Initialize models
log_reg = LogisticRegression(random_state=35, max_iter=1000,
    ↪class_weight='balanced',
                                solver='lbfgs')
rf_clf = RandomForestClassifier(random_state=35, class_weight='balanced',
                                n_estimators=100, max_depth=10)
svm_clf = SVC(kernel='rbf', # RBF kernel for non-linear classification
               C=1.0,       # Regularization parameter
               gamma='scale', # Kernel coefficient
               class_weight='balanced',
               probability=True,
               random_state=35)

# Dictionary to store results
models_results = {}

# Train and evaluate all models
models_results['Logistic Regression'] = train_and_evaluate_model(
    log_reg, X_train, X_test, y_train, y_test, "Logistic Regression"
)

models_results['Random Forest'] = train_and_evaluate_model(
    rf_clf, X_train, X_test, y_train, y_test, "Random Forest"
)

models_results['SVM'] = train_and_evaluate_model(
    svm_clf, X_train, X_test, y_train, y_test, "SVM", needs_scaling=True # SVM
    ↪needs scaled features
)

# Compare models
compare_models(models_results)

# Feature importance for Random Forest and coefficients for Logistic Regression
if hasattr(X_train, 'columns'):
    # Random Forest feature importance
    rf_importance = pd.DataFrame({
        'feature': X_train.columns,
        'importance': rf_clf.feature_importances_
    }).sort_values('importance', ascending=False)

    # Logistic Regression coefficients (absolute values)
    lr_importance = pd.DataFrame({
        'feature': X_train.columns,
        'importance': np.abs(log_reg.coef_[0])
    }).sort_values('importance', ascending=False)

```

```

# Plot feature importance
plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
sns.barplot(data=rf_importance, x='importance', y='feature')
plt.title('Random Forest Feature Importance')

plt.subplot(1, 2, 2)
sns.barplot(data=lr_importance, x='importance', y='feature')
plt.title('Logistic Regression Feature Importance')

plt.tight_layout()
plt.show()

# Calculate and display model comparison statistics
comparison_stats = pd.DataFrame({
    'Model': models_results.keys(),
    'Test Accuracy': [results['accuracy'] for results in models_results.
↳values()],
    'CV Mean Accuracy': [results['cv_scores'].mean() for results in
↳models_results.values()],
    'CV Std': [results['cv_scores'].std() for results in models_results.
↳values()]
})

print("\nModel Comparison Summary:")
print("-" * 80)
print(comparison_stats.to_string(index=False))

```

Logistic Regression Results:

-----  
Test Set Accuracy: 0.8280


K-fold Cross Validation Scores:

Mean CV Accuracy: 0.8412 (+/- 0.0372)

Individual fold scores: ['0.8167', '0.8247', '0.8526', '0.8680', '0.8440']

Classification Report:

	precision	recall	f1-score	support
0	0.12	0.23	0.16	22
1	0.94	0.87	0.90	292
accuracy			0.83	314
macro avg	0.53	0.55	0.53	314
weighted avg	0.88	0.83	0.85	314



Random Forest Results:

-----  
Test Set Accuracy: 0.9299

K-fold Cross Validation Scores:

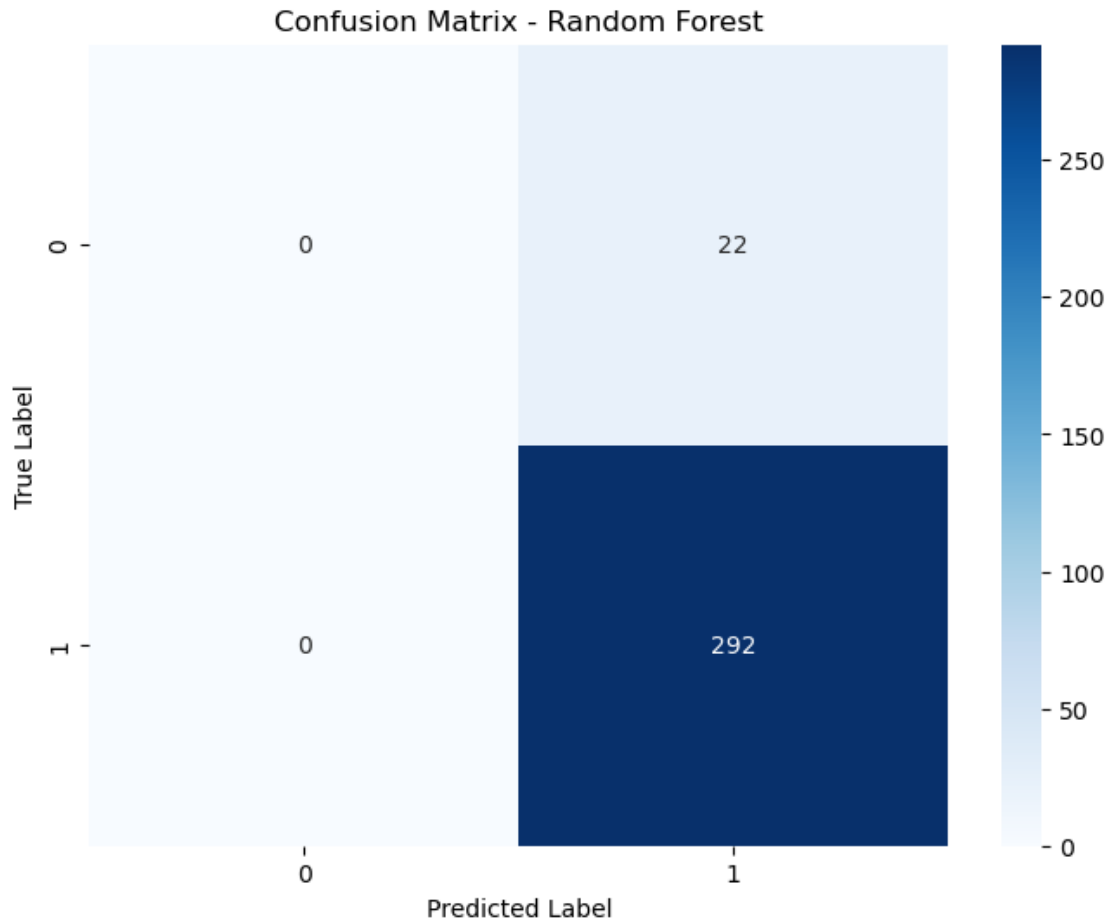
Mean CV Accuracy: 0.9345 (+/- 0.0374)

Individual fold scores: ['0.9163', '0.9522', '0.9522', '0.9080', '0.9440']

Classification Report:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	22
1	0.93	1.00	0.96	292
accuracy			0.93	314
macro avg	0.46	0.50	0.48	314

weighted avg      0.86      0.93      0.90      314



SVM Results:

-----  
Test Set Accuracy: 0.9204

K-fold Cross Validation Scores:


Mean CV Accuracy: 0.9290 (+/- 0.0348)

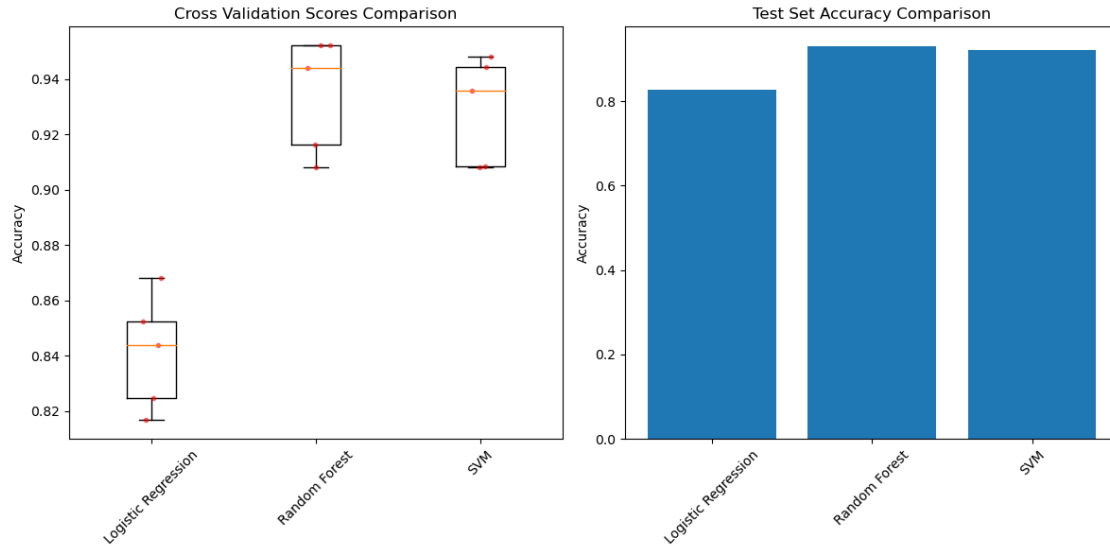
Individual fold scores: ['0.9084', '0.9442', '0.9482', '0.9080', '0.9360']

Classification Report:

	precision	recall	f1-score	support
0	0.20	0.05	0.07	22
1	0.93	0.99	0.96	292
accuracy			0.92	314

macro avg	0.57	0.52	0.52	314
weighted avg	0.88	0.92	0.90	314





Model Comparison Summary:

Model	Test Accuracy	CV Mean Accuracy	CV Std
Logistic Regression	0.828025	0.841205	0.018583
Random Forest	0.929936	0.934543	0.018700
SVM	0.920382	0.928959	0.017415

## 2.1 Artificial Neural Network

We will try to optimize the neural network parameters for the same problem. We will then evaluate how well the ANN model predicts the pass or fail criteria.

```
[209]: import tensorflow as tf
from tensorflow import keras
from keras import callbacks
```

Since the Dataset has 436 input features, we will also look into whether apply PCA is beneficial for generalization of the model as well as to improve computing time.

1. Without PCA, feed all the 436 features into ANN model
2. Apply PCA, and optimize parameters for the ANN model

```
[210]: import numpy as np
from sklearn.model_selection import KFold
from sklearn.decomposition import PCA

def create_model():
    # Initialize the ANN model
    model = tf.keras.models.Sequential()
```

```

# Input Layer
model.add(tf.keras.layers.Dense(units=436, activation='relu'))

# Hidden Layers with Dropout and Regularization
model.add(tf.keras.layers.Dense(700, activation='relu',
↳kernel_regularizer=tf.keras.regularizers.l2(0.01)))
model.add(tf.keras.layers.Dropout(0.2))
model.add(tf.keras.layers.Dense(400, activation='relu',
↳kernel_regularizer=tf.keras.regularizers.l2(0.01)))
model.add(tf.keras.layers.Dropout(0.3))

# Output Layer
model.add(tf.keras.layers.Dense(1, activation='sigmoid')) # Change to
↳softmax for multi-class

# Compile the model
model.compile(optimizer='adam', loss='binary_crossentropy',
↳metrics=['accuracy']) # Change loss for multi-class
return model

# Ensure X and y are numpy arrays
X_train = np.array(X_train)
y_train = np.array(y_train)

# # Fit PCA
# pca = PCA(n_components=0.95) # Retain 95% of the variance
# X_pca = pca.fit_transform(X)

# Initialize K-fold cross-validation
n_splits = 5
kfold = KFold(n_splits=n_splits, shuffle=True, random_state=42)

# Lists to store metrics
fold accuracies = []
fold losses = []
histories = []
all_predictions = []
all_true_values = []

# K-fold cross validation
for fold, (train_idx, val_idx) in enumerate(kfold.split(X_train)):
    print(f'\nFold {fold + 1}/{n_splits}')

    # Split data
    X_train_fold = X_train[train_idx]
    y_train_fold = y[train_idx]

```



```

X_val_fold = X_train[val_idx]
y_val_fold = y[val_idx]

# Create and compile model
model = create_model()

# Callbacks
lr_scheduler = callbacks.ReduceLROnPlateau(
    monitor='val_loss',
    factor=0.5,
    patience=5
)
early_stopping = callbacks.EarlyStopping(
    monitor='val_loss',
    patience=10,
    restore_best_weights=True
)

# Train model
history = model.fit(
    X_train_fold,
    y_train_fold,
    batch_size=256,
    epochs=100,
    validation_data=(X_val_fold, y_val_fold),
    callbacks=[lr_scheduler, early_stopping],
    verbose=0
)

# Evaluate model
loss, accuracy = model.evaluate(X_val_fold, y_val_fold, verbose=0)
fold accuracies.append(accuracy)
fold_losses.append(loss)
histories.append(history.history)

# Generate predictions
y_pred = model.predict(X_val_fold, verbose=0)
y_pred_binary = (y_pred > 0.5).astype(int).reshape(-1) # Ensure 1D array

# Store predictions and true values
all_predictions.extend(y_pred_binary)
all_true_values.extend(y_val_fold)

# Calculate metrics for this fold
conf_matrix = confusion_matrix(y_val_fold, y_pred_binary)

print(f'\nFold {fold + 1} Results:')

```

```

print(f'Validation Loss: {loss:.4f}')
print(f'Validation Accuracy: {accuracy:.4f}')
print('\nConfusion Matrix:')

# Plot Confusion Matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=['Fail', 'Pass'], yticklabels=['Fail', 'Pass'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()

print('\nClassification Report:')
print(classification_report(y_val_fold, y_pred_binary,
    ↪target_names=['Fail', 'Pass']))

# Convert to numpy arrays for final calculations
all_predictions = np.array(all_predictions)
all_true_values = np.array(all_true_values)

# Print overall results
print('\nOverall K-fold Cross Validation Results:')
print('-----')
print(f'Mean Accuracy: {np.mean(fold accuracies):.4f} (±{np.
    ↪std(fold accuracies):.4f})')
print(f'Mean Loss: {np.mean(fold losses):.4f} (±{np.std(fold losses):.4f})')

# Overall confusion matrix and classification report
print('\nOverall Confusion Matrix:')
print(confusion_matrix(all_true_values, all_predictions))

```


Fold 1/5

Fold 1 Results:

Validation Loss: 0.4714

Validation Accuracy: 0.9283

Confusion Matrix:



Classification Report:

	precision	recall	f1-score	support
Fail	0.00	0.00	0.00	18
Pass	0.93	1.00	0.96	233
accuracy			0.93	251
macro avg	0.46	0.50	0.48	251
weighted avg	0.86	0.93	0.89	251

Fold 2/5

```
c:\Users\kylea\anaconda3\lib\site-
packages\sklearn\metrics\_classification.py:1531: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
c:\Users\kylea\anaconda3\lib\site-
```

```
packages\sklearn\metrics\_classification.py:1531: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
c:\Users\kylea\anaconda3\lib\site-
```

```
packages\sklearn\metrics\_classification.py:1531: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
```


```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

Fold 2 Results:

Validation Loss: 0.3533

Validation Accuracy: 0.9482

Confusion Matrix:



Classification Report:

precision	recall	f1-score	support
-----------	--------	----------	---------

Fail	0.00	0.00	0.00	12
Pass	0.95	1.00	0.97	239
accuracy			0.95	251
macro avg	0.48	0.50	0.49	251
weighted avg	0.91	0.95	0.93	251


Fold 3/5

Fold 3 Results:

Validation Loss: 0.5179

Validation Accuracy: 0.9203

Confusion Matrix:



Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

Fail	0.00	0.00	0.00	18
Pass	0.93	0.99	0.96	233
accuracy			0.92	251
macro avg	0.46	0.50	0.48	251
weighted avg	0.86	0.92	0.89	251

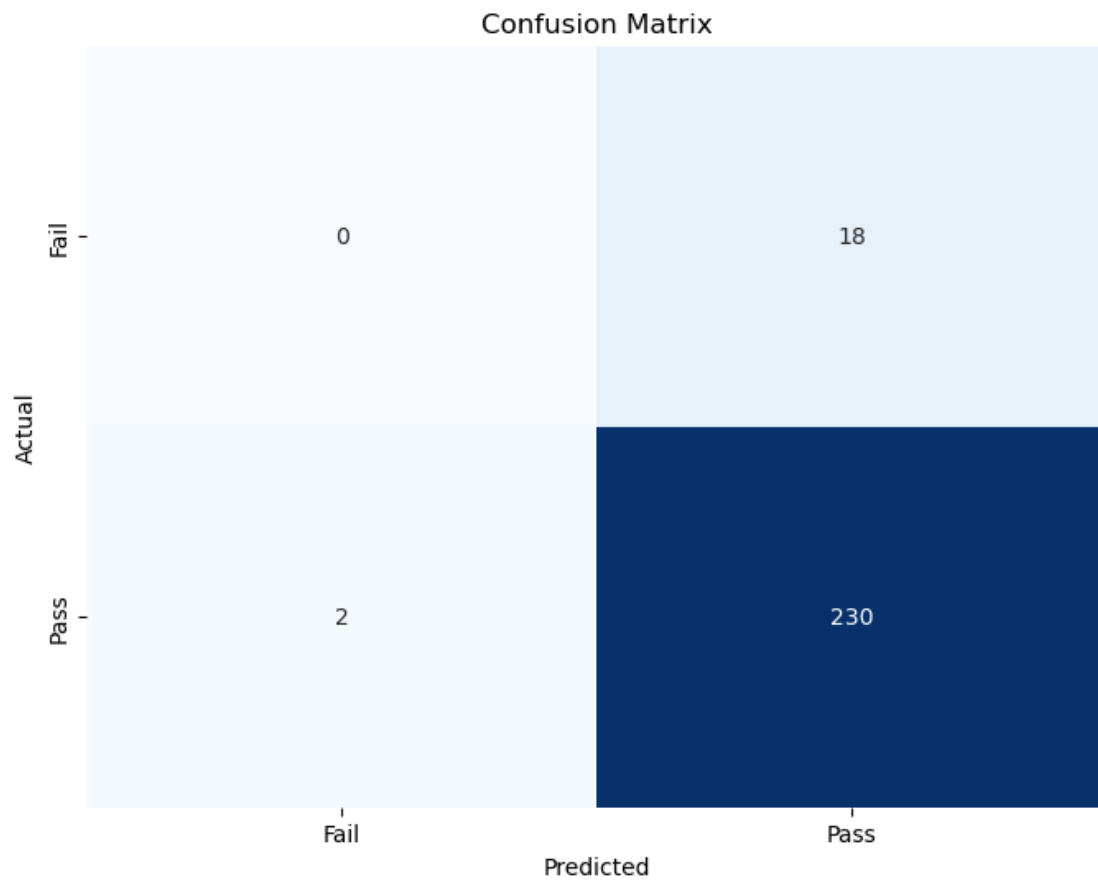
Fold 4/5

Fold 4 Results:

Validation Loss: 0.4936

Validation Accuracy: 0.9200

Confusion Matrix:



Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

Fail	0.00	0.00	0.00	18
Pass	0.93	0.99	0.96	232
accuracy			0.92	250
macro avg	0.46	0.50	0.48	250
weighted avg	0.86	0.92	0.89	250


Fold 5/5

Fold 5 Results:

Validation Loss: 0.6383

Validation Accuracy: 0.9160

Confusion Matrix:



Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

Fail	0.00	0.00	0.00	21
Pass	0.92	1.00	0.96	229
accuracy			0.92	250
macro avg	0.46	0.50	0.48	250
weighted avg	0.84	0.92	0.88	250

Overall K-fold Cross Validation Results:

-----  
Mean Accuracy: 0.9266 ( $\pm 0.0115$ )  
Mean Loss: 0.4949 ( $\pm 0.0913$ )

Overall Confusion Matrix:

```
[[ 0  87]
 [ 5 1161]]
```

```
c:\Users\kylea\anaconda3\lib\site-
packages\sklearn\metrics\_classification.py:1531: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
c:\Users\kylea\anaconda3\lib\site-
packages\sklearn\metrics\_classification.py:1531: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
c:\Users\kylea\anaconda3\lib\site-
packages\sklearn\metrics\_classification.py:1531: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

## 2.2 Implement PCA - ANN

```
[211]: from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler

# You can also concatenate your train and test data if needed for PCA fitting
# X = np.vstack((X_train, X_test))

# Standardize the data (important for PCA)
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_train)

# Initialize PCA
pca = PCA(n_components=2) # Adjust the number of components as needed
```



```

# Fit PCA on the scaled data
X_pca = pca.fit_transform(X_scaled)


# Create a DataFrame for PCA results
pca_df = pd.DataFrame(data=X_pca, columns=['Principal Component 1', 'Principal_
Component 2'])

# Explained variance ratio
explained_variance = pca.explained_variance_ratio_
print(f'Explained variance ratio: {explained_variance}')

# Plotting the PCA results
plt.figure(figsize=(8, 6))
plt.scatter(pca_df['Principal Component 1'], pca_df['Principal Component 2'],
alpha=0.7)
plt.title('PCA Result')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.grid()
plt.show()

```

Explained variance ratio: [0.06044871 0.03900913]




```
[212]: # Fit PCA again to get explained variance ratio
pca = PCA()
pca.fit(X_scaled)

# Plot the explained variance
plt.figure(figsize=(10, 6))
plt.plot(range(1, len(pca.explained_variance_ratio_) + 1),
         np.cumsum(pca.explained_variance_ratio_),
         marker='o', linestyle='--')
plt.title('Cumulative Explained Variance by PCA Components')
plt.xlabel('Number of Components')
plt.ylabel('Cumulative Explained Variance')
plt.axhline(y=0.95, color='r', linestyle='--') # Line for 95% variance
plt.axvline(x=np.argmax(np.cumsum(pca.explained_variance_ratio_) >= 0.95) + 1,
           color='g', linestyle='--') # Line for the number of components needed for
           95%
plt.text(1, 0.9, '95% Variance', color='red')
```

```
plt.text(np.argmax(np.cumsum(pca.explained_variance_ratio_) >= 0.95) + 1, 0.8,
        f'Components: {np.argmax(np.cumsum(pca.explained_variance_ratio_) >= 0.95) + 1}', color='green')
plt.grid()
plt.show()

# Print the number of components used in PCA
n_components = np.argmax(np.cumsum(pca.explained_variance_ratio_) >= 0.95) + 1
print(f"Number of components retained to achieve 95% variance: {n_components}")
```



Number of components retained to achieve 95% variance: 157

```
[213]: import numpy as np
from sklearn.model_selection import KFold
from sklearn.decomposition import PCA

def create_model():
    # Initialize the ANN model
    model = tf.keras.models.Sequential()

    # Input Layer
    model.add(tf.keras.layers.Dense(units=157, activation='relu'))

    # Hidden Layers with Dropout
```

```

    model.add(tf.keras.layers.Dense(252, activation='relu',
↪kernel_regularizer=tf.keras.regularizers.l2(0.01)))
    model.add(tf.keras.layers.Dropout(0.2))
    model.add(tf.keras.layers.Dense(150, activation='relu',
↪kernel_regularizer=tf.keras.regularizers.l2(0.01)))
    model.add(tf.keras.layers.Dropout(0.3))

    # Output Layer
    model.add(tf.keras.layers.Dense(1, activation='sigmoid')) # Change to
↪softmax for multi-class

    # Compile the model
    model.compile(optimizer='adam', loss='binary_crossentropy',
↪metrics=['accuracy']) # Change loss for multi-class
    return model

# Ensure X and y are numpy arrays
X = np.array(X)
y = np.array(y)

# Fit PCA
pca = PCA(n_components=0.95) # Retain 95% of the variance
X_pca = pca.fit_transform(X)

# Initialize K-fold cross-validation
n_splits = 5
kfold = KFold(n_splits=n_splits, shuffle=True, random_state=42)

# Lists to store metrics
fold_accuracies = []
fold_losses = []
histories = []
all_predictions = []
all_true_values = []

# K-fold cross validation
for fold, (train_idx, val_idx) in enumerate(kfold.split(X_pca)):
    print(f'\nFold {fold + 1}/{n_splits}')

    # Split data
    X_train_fold = X_pca[train_idx]
    y_train_fold = y[train_idx]
    X_val_fold = X_pca[val_idx]
    y_val_fold = y[val_idx]

    # Create and compile model
    model = create_model()

```

```

# Callbacks
lr_scheduler = callbacks.ReduceLROnPlateau(
    monitor='val_loss',
    factor=0.5,
    patience=5
)
early_stopping = callbacks.EarlyStopping(
    monitor='val_loss',
    patience=10,
    restore_best_weights=True
)

# Train model
history = model.fit(
    X_train_fold,
    y_train_fold,
    batch_size=256,
    epochs=100,
    validation_data=(X_val_fold, y_val_fold),
    callbacks=[lr_scheduler, early_stopping],
    verbose=0
)

# Evaluate model
loss, accuracy = model.evaluate(X_val_fold, y_val_fold, verbose=0)
fold accuracies.append(accuracy)
fold_losses.append(loss)
histories.append(history.history)

# Generate predictions
y_pred = model.predict(X_val_fold, verbose=0)
y_pred_binary = (y_pred > 0.5).astype(int).reshape(-1) # Ensure 1D array

# Store predictions and true values
all_predictions.extend(y_pred_binary)
all_true_values.extend(y_val_fold)

# Calculate metrics for this fold
conf_matrix = confusion_matrix(y_val_fold, y_pred_binary)

print(f'\nFold {fold + 1} Results:')
print(f'Validation Loss: {loss:.4f}')
print(f'Validation Accuracy: {accuracy:.4f}')
print(f'\nConfusion Matrix:')

# Plot Confusion Matrix

```

```

plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=['Fail', 'Pass'], yticklabels=['Fail', 'Pass'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()

print('\nClassification Report:')
print(classification_report(y_val_fold, y_pred_binary,
    ↪target_names=['Fail', 'Pass']))

# Convert to numpy arrays for final calculations
all_predictions = np.array(all_predictions)
all_true_values = np.array(all_true_values)

# Print overall results
print('\nOverall K-fold Cross Validation Results:')
print('-----')
print(f'Mean Accuracy: {np.mean(fold accuracies):.4f} (±{np.
    ↪std(fold accuracies):.4f})')
print(f'Mean Loss: {np.mean(fold losses):.4f} (±{np.std(fold losses):.4f})')

# Overall confusion matrix and classification report
print('\nOverall Confusion Matrix:')
print(confusion_matrix(all_true_values, all_predictions))

```

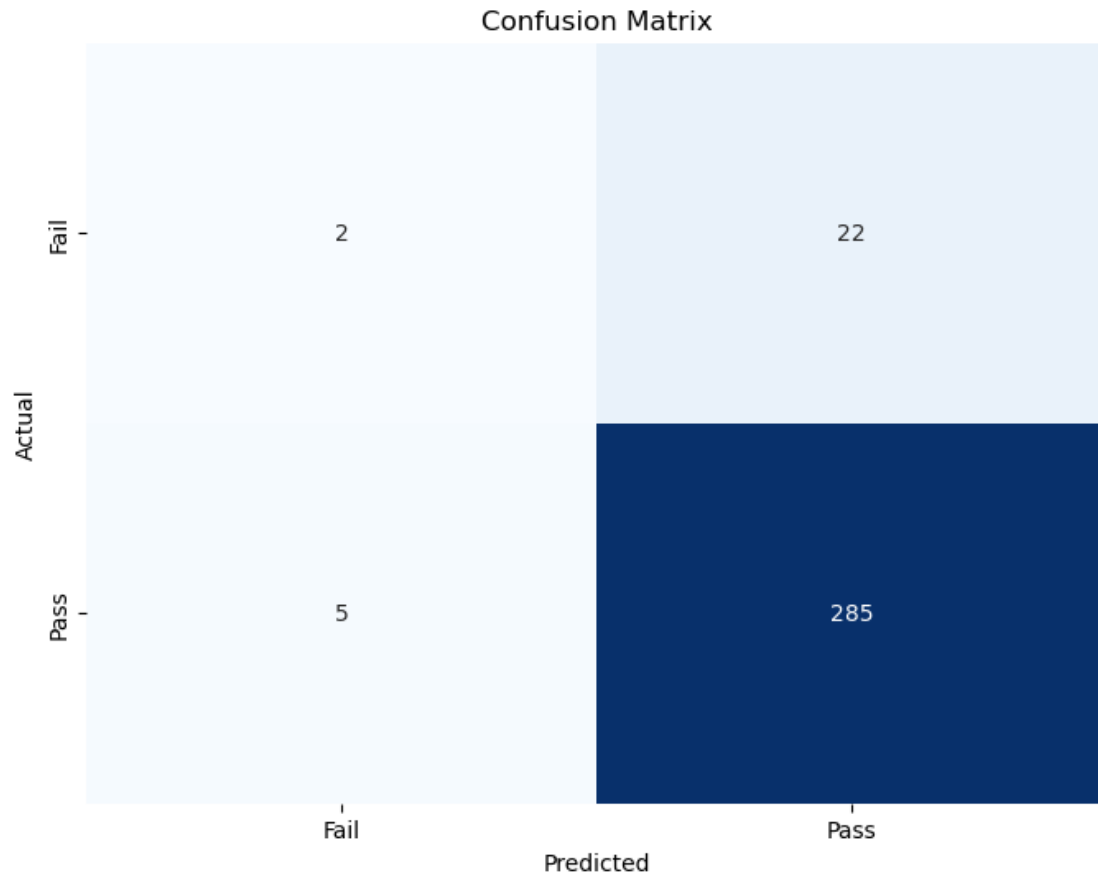
Fold 1/5

Fold 1 Results:

Validation Loss: 0.4326

Validation Accuracy: 0.9140

Confusion Matrix:



Classification Report:

	precision	recall	f1-score	support
Fail	0.29	0.08	0.13	24
Pass	0.93	0.98	0.95	290
accuracy			0.91	314
macro avg	0.61	0.53	0.54	314
weighted avg	0.88	0.91	0.89	314


Fold 2/5

Fold 2 Results:

Validation Loss: 0.2815

Validation Accuracy: 0.9459

Confusion Matrix:



Classification Report:

	precision	recall	f1-score	support
Fail	0.00	0.00	0.00	17
Pass	0.95	1.00	0.97	297
accuracy			0.95	314
macro avg	0.47	0.50	0.49	314
weighted avg	0.89	0.95	0.92	314

Fold 3/5

```
c:\Users\kylea\anaconda3\lib\site-
packages\sklearn\metrics\_classification.py:1531: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
c:\Users\kylea\anaconda3\lib\site-
```



```
packages\sklearn\metrics\_classification.py:1531: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```


```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
c:\Users\kylea\anaconda3\lib\site-
packages\sklearn\metrics\_classification.py:1531: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

Fold 3 Results:

Validation Loss: 0.4124

Validation Accuracy: 0.9233

Confusion Matrix:



Classification Report:

precision	recall	f1-score	support
-----------	--------	----------	---------

Fail	0.00	0.00	0.00	23
Pass	0.93	1.00	0.96	290
accuracy			0.92	313
macro avg	0.46	0.50	0.48	313
weighted avg	0.86	0.92	0.89	313

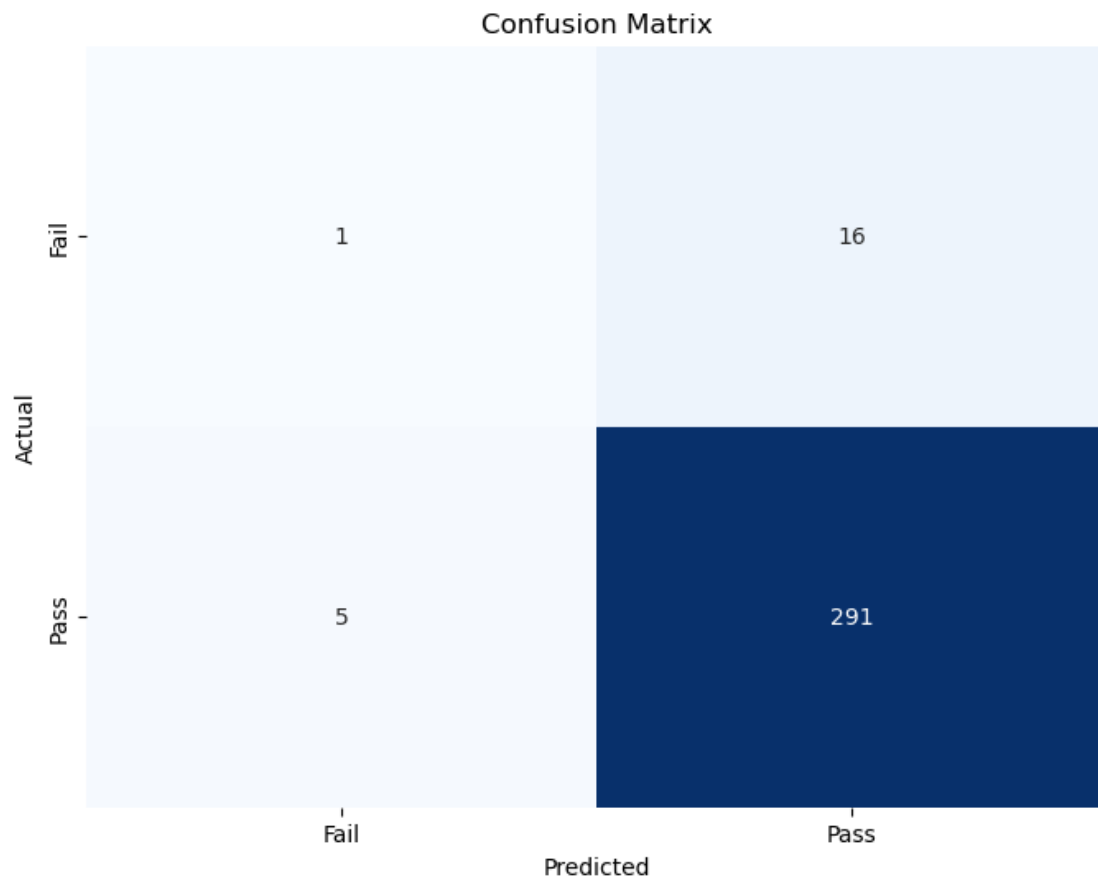
Fold 4/5

Fold 4 Results:

Validation Loss: 0.3627

Validation Accuracy: 0.9329

Confusion Matrix:



Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

Fail	0.17	0.06	0.09	17
Pass	0.95	0.98	0.97	296
accuracy			0.93	313
macro avg	0.56	0.52	0.53	313
weighted avg	0.91	0.93	0.92	313


Fold 5/5

Fold 5 Results:

Validation Loss: 0.4505

Validation Accuracy: 0.9297

Confusion Matrix:



Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

Fail	1.00	0.04	0.08	23
Pass	0.93	1.00	0.96	290
accuracy			0.93	313
macro avg	0.96	0.52	0.52	313
weighted avg	0.93	0.93	0.90	313

Overall K-fold Cross Validation Results:

-----

Mean Accuracy: 0.9292 ( $\pm 0.0106$ )

Mean Loss: 0.3879 ( $\pm 0.0608$ )

Overall Confusion Matrix:

```
[[ 4 100]
 [ 11 1452]]
```

### 3 Summary

## 4 Neural Network Architecture and Model Comparison Analysis

### 4.1 Model Architectures

#### 4.1.1 ANN without PCA

- First Hidden Layer: 700 nodes
- Second Hidden Layer: 400 nodes
- Dropout Rate: 0.3
- Test Accuracy: 92.66%
- Loss Function: 0.50

#### 4.1.2 ANN with PCA (95% variance retained)

- First Hidden Layer: 252 nodes
- Second Hidden Layer: 150 nodes
- Test Accuracy: 92.92%
- Loss Function: 0.39

#### 4.1.3 Traditional Classification Models

Model	Test Accuracy	CV Mean Accuracy	CV Std
Logistic Regression	82.80%	84.13%	0.018583
Random Forest	92.99%	93.45%	0.018700
SVM	92.04%	92.90%	0.017415

## 4.2 Conclusion

The implementation and comparison of various machine learning models for this classification task has revealed several interesting insights:

### 4.2.1 1. PCA Implementation Impact

- The ANN with PCA slightly outperformed the non-PCA version (92.92% vs 92.66%)
- More significantly, PCA reduced the loss function from 0.50 to 0.39
- The dimensional reduction allowed for a more efficient architecture (252/150 nodes vs 700/400 nodes)
- The maintained performance despite reduced complexity suggests effective feature extraction

### 4.2.2 2. Model Architecture Optimization

- The implementation of dropout (0.3) and regularization successfully addressed overfitting
- The reduced architecture with PCA achieved comparable results with significantly fewer parameters
- The more compact architecture likely contributed to the reduced loss function

### 4.2.3 3. Comparative Performance

- Random Forest emerged as the top performer with 92.99% test accuracy and 93.45% CV mean accuracy
- Both ANN implementations (with and without PCA) performed competitively with Random Forest
- SVM showed strong performance (92.04% test accuracy, 92.90% CV mean accuracy)
- Logistic Regression significantly underperformed compared to other models

### 4.2.4 4. Model Stability

- All advanced models (Random Forest, SVM, and both ANNs) achieved >92% accuracy
- Cross-validation standard deviations were consistently low (~0.018), indicating stable performance
- The similar performance across different architectural approaches suggests robust feature relationships

### 4.2.5 5. Efficiency Considerations

- PCA implementation provided a more efficient solution without sacrificing performance
- The reduced architecture with PCA suggests potential benefits for computational efficiency
- The lower loss function with PCA indicates better model optimization

## 4.3 Recommendations

1. For production deployment, both the Random Forest and ANN with PCA would be strong candidates:
  - Random Forest for its slightly higher accuracy and proven stability
  - ANN with PCA for its efficient architecture and lower loss function
2. Future optimization could focus on:

- Fine-tuning the PCA variance retention percentage
- Exploring ensemble methods combining the strengths of multiple models
- Further architecture optimization for the neural networks

The results demonstrate that while traditional methods like Random Forest remain competitive, modern approaches with dimensional reduction and neural networks can achieve comparable results with potentially more efficient architectures.